

Meeting 11/29/17

Wednesday, November 29, 2017 11:00 AM

Agenda:

Attendees: Everyone

Compression didn't apply for small data

Entered rows with the same item to see what it did with the token (cardinality of 1)
Also interwove rows (ABAB) and it had 4 tokens

1 2 3 4
AA AA AA AA
BB BB BB BB

Token Table

- 1) AA
- 2) BB
- 3) 1.1.1.1
- 4) 2.2.2.2

Table

- 1) 3
- 2) 4
- 3) 3
- 4) 4

Bindmp stored in raw data or pulled from header / somewhere else?

Experiment

- Run size (length) of varchar
 - o Run cardinality of the varchar
 - How much space does an object take up normalized / de-normalized as a function of size and cardinality
- Column ordering / reordering
 - o When and why does it reorder columns
 - Length and cardinality

Results: Columns are ordered as inserted without ORDER BY
When ORDER BY is used, that's how it's ordered

ORDER BY 5 4 3 2 1 0 results in 2 3 4 5 0 1
0 = 5
1 = 5
2 = 16
3 = 7
4 = 8
5 = 32

Block cardinality of 1 doesn't matter too much (maybe)

Try order by vs let it do its own thing

Order by different columns

Make everything the same size, no var char 5-9

For Next Week:

- Figure out when and why columns are reordered
- (draw notes from notebook)
- Continue deconstructing the block BINDMP and TOKENS